

# Introducing Riak and Ripple

Sean Cribbs

system “whoami”

# system “whoami”

- BCS, three years as a musician

# system “whoami”

- BCS, three years as a musician
- Started with Ruby in early 2006

# system “whoami”

- BCS, three years as a musician
- Started with Ruby in early 2006
- Rails freelancer 2007-2010

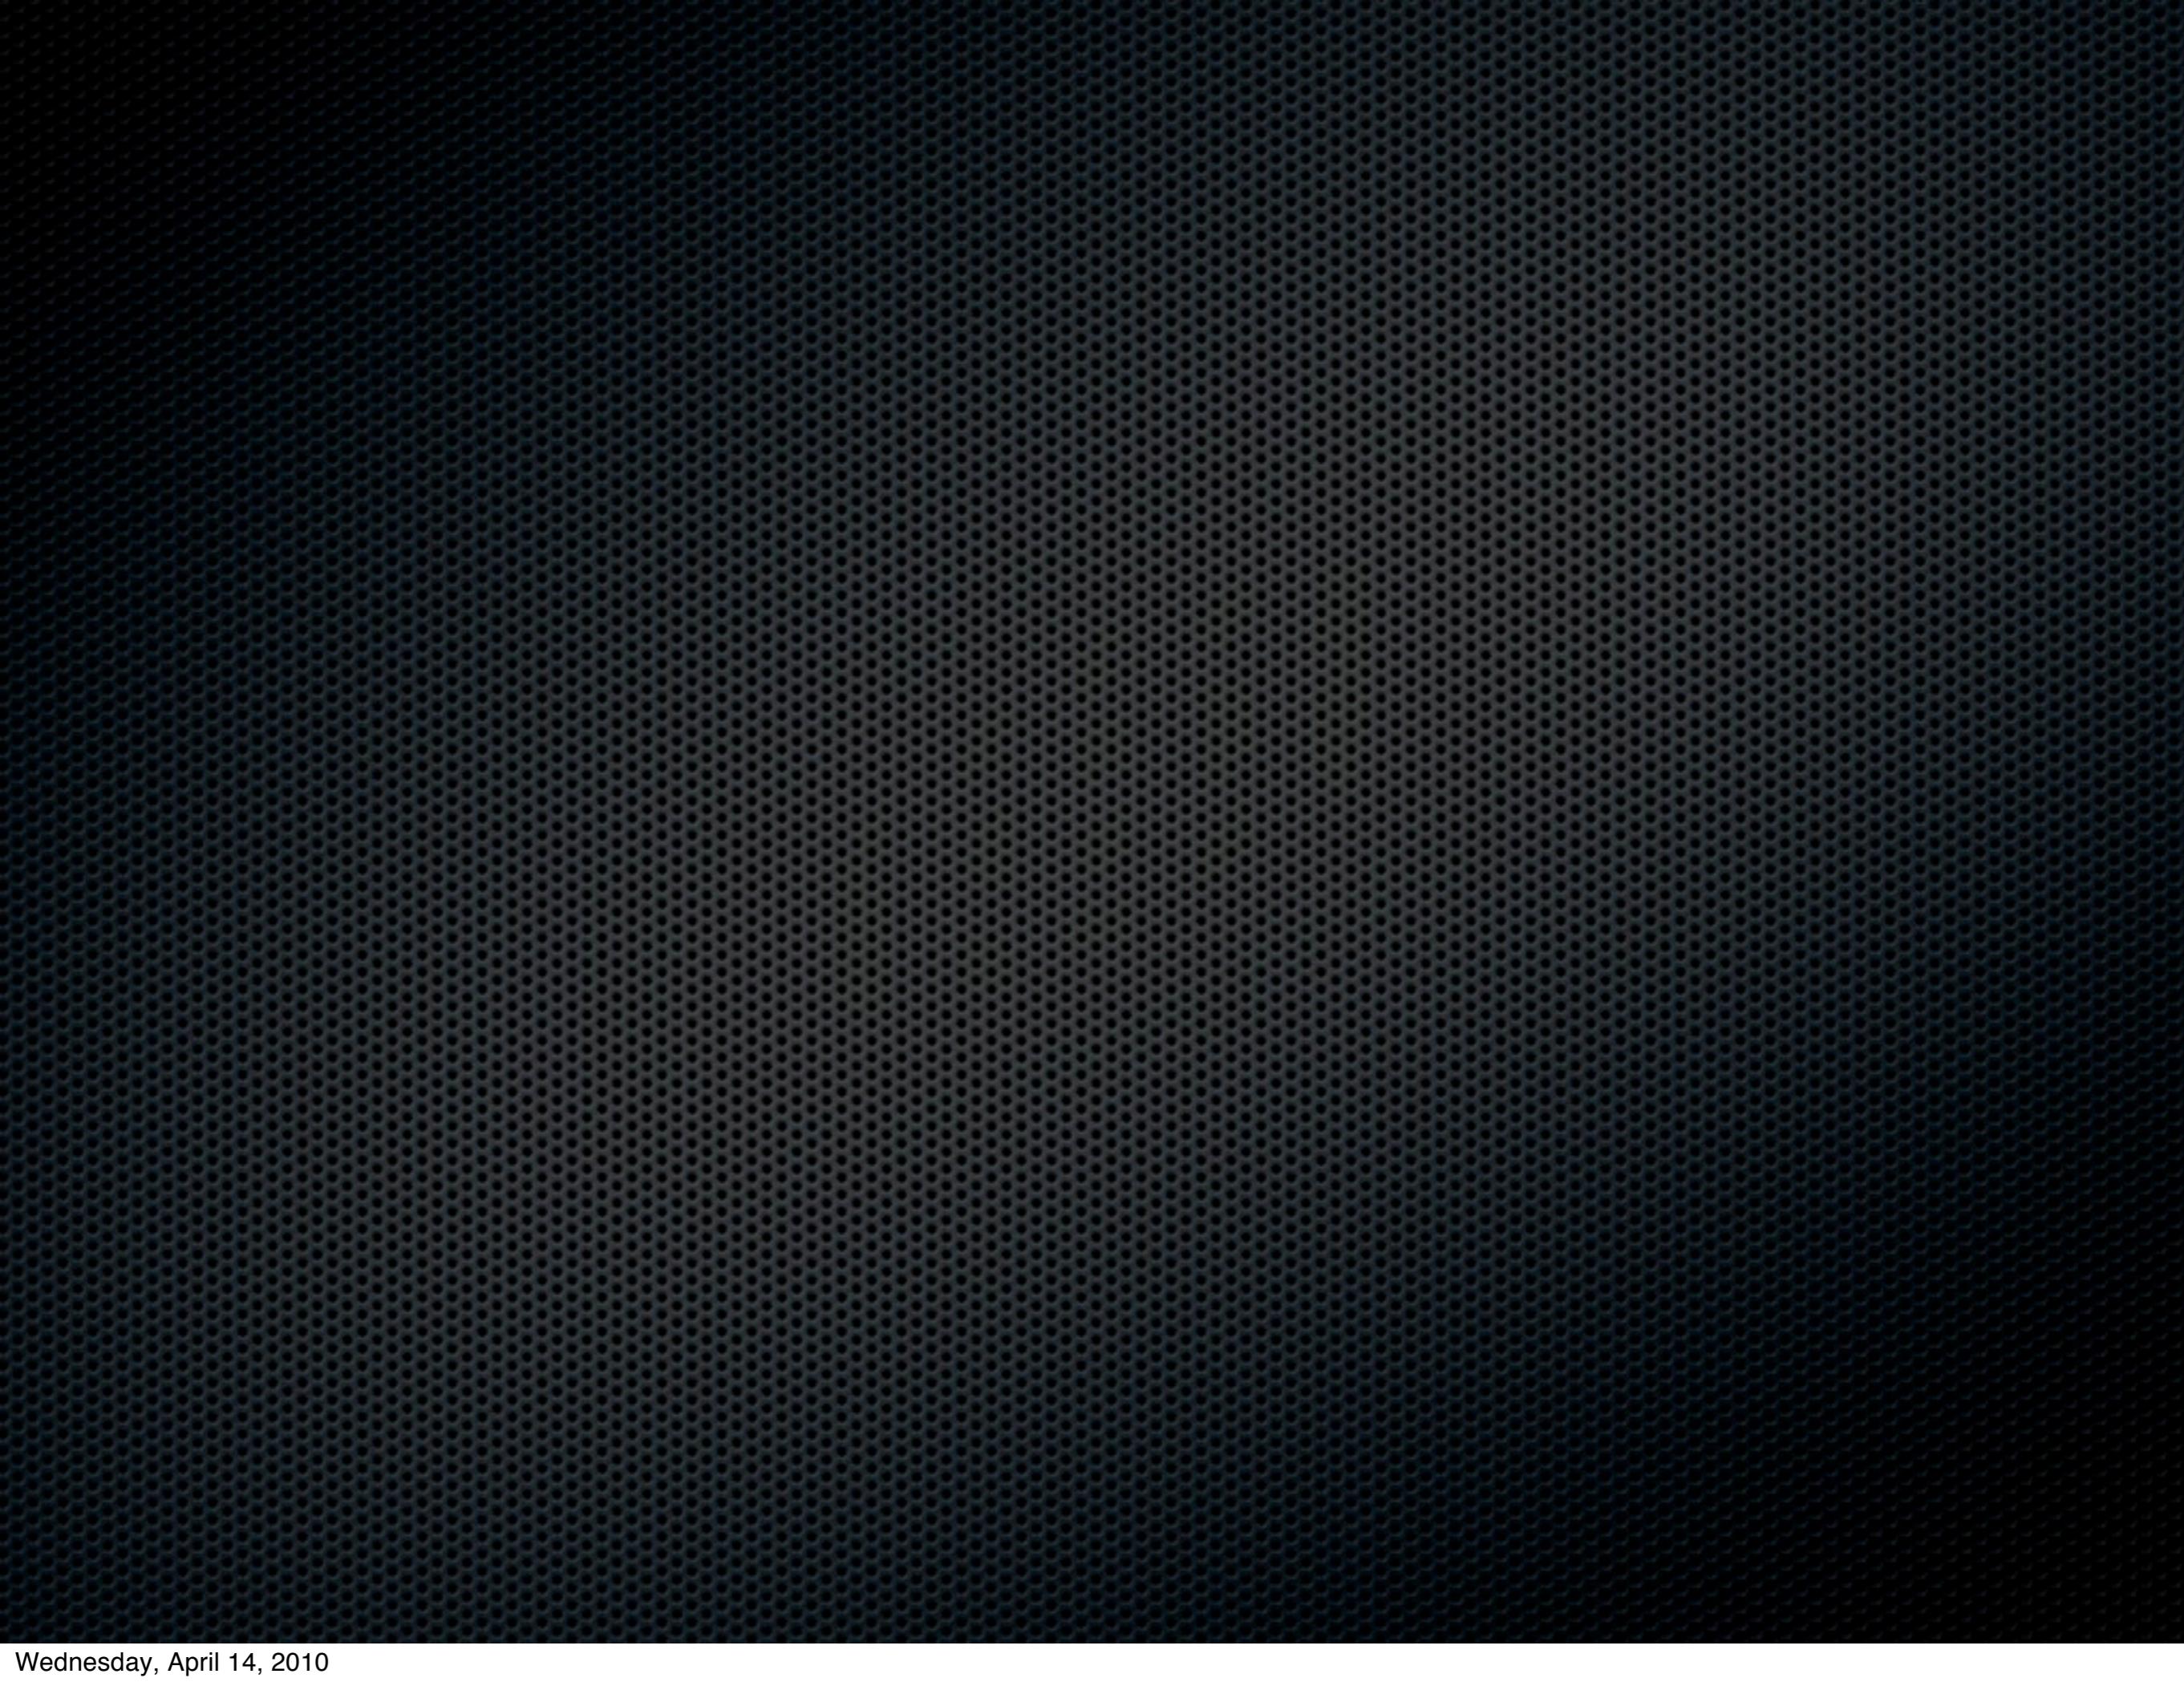
# system “whoami”

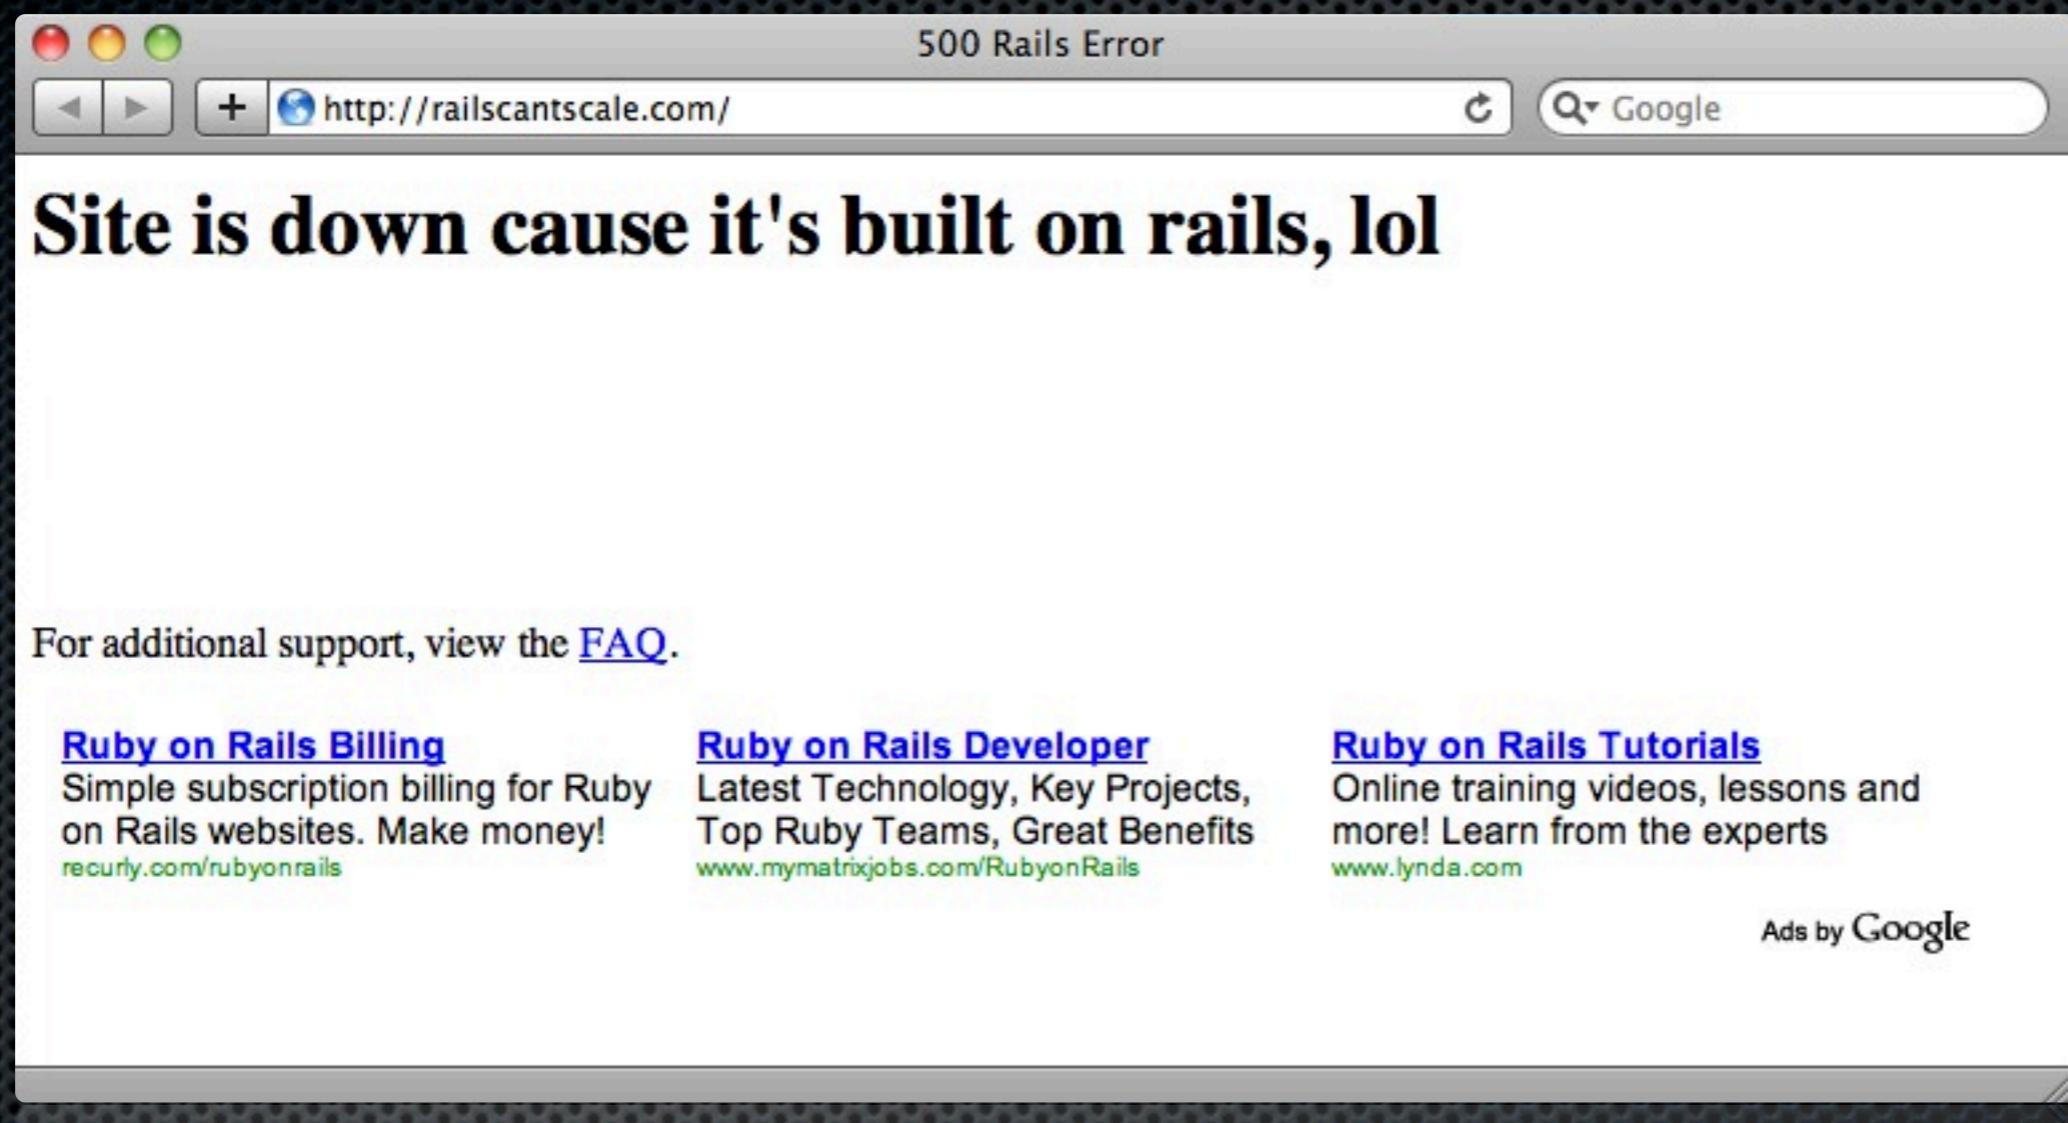
- BCS, three years as a musician
- Started with Ruby in early 2006
- Rails freelancer 2007-2010
- Radiant CMS

# system “whoami”

- BCS, three years as a musician
- Started with Ruby in early 2006
- Rails freelancer 2007-2010
- Radiant CMS
- March 2010 - Developer Advocate







# Rails Can't Scale

# Rails Can't Scale

what does it mean?

# Ruby ~~Rails~~ Can't Scale

what does it mean?

# MySQL ~~Ruby~~ Rails Can't Scale

what does it mean?

*insert straw man here*

~~MySQL~~ ←  
~~Ruby~~ ~~Rails~~ Can't Scale

what does it mean?

# Mythbusting “scalability”

# Mythbusting “scalability”

- Scalability is not a yes/no question

# Mythbusting “scalability”

- Scalability is not a yes/no question
- It's a ratio of **benefit** to **cost**

# Mythbusting “scalability”

- Scalability is not a yes/no question
- It's a ratio of **benefit** to **cost**
- **Benefits:** low latency, throughput, uptime, concurrency, reliability

# Mythbusting “scalability”

- Scalability is not a yes/no question
- It's a ratio of **benefit** to **cost**
- **Benefits:** low latency, throughput, uptime, concurrency, reliability
- **Costs:** CPU, RAM, disk, bandwidth, power, hw/sw

# Mythbusting “scalability”

- Scalability is not a yes/no question
- It's a ratio of **benefit** to **cost**
  - **Benefits:** low latency, throughput, uptime, concurrency, reliability
  - **Costs:** CPU, RAM, disk, bandwidth, power, hw/sw
- **scalability = bang for your buck (ROI)**

# Scaling goes both ways

# Scaling goes both ways



- **Up** - high traffic, “Big Data”, multi-datacenter

# Scaling goes both ways



- **Up** - high traffic, “Big Data”, multi-datacenter
- **Down** - laptop, set-top, mobile



# How do you scale?

# How do you scale?

- Vertically - get bigger, expensive machines

# How do you scale?

- Vertically - get bigger, expensive machines
- Horizontally - get more commodity machines

# What is Riak?

(the horizontal scaling bits)

# What is Riak?

(the horizontal scaling bits)

- Based on Amazon's Dynamo (2007)

# What is Riak?

(the horizontal scaling bits)

- Based on Amazon's Dynamo (2007)
  - key-value storage

# What is Riak?

(the horizontal scaling bits)

- Based on Amazon's Dynamo (2007)
  - key-value storage
  - masterless, peer-to-peer replication

# What is Riak?

(the horizontal scaling bits)

- Based on Amazon's Dynamo (2007)
  - key-value storage
  - masterless, peer-to-peer replication
  - consistent hashing

# What is Riak?

(the horizontal scaling bits)

- Based on Amazon's Dynamo (2007)
  - key-value storage
  - masterless, peer-to-peer replication
  - consistent hashing
  - eventual consistency

# What is Riak?

(the horizontal scaling bits)

- Based on Amazon's Dynamo (2007)
  - key-value storage
  - masterless, peer-to-peer replication
  - consistent hashing
  - eventual consistency
  - failover - quorums, hinted handoff

*“O/RM is the Vietnam of Computer Science.”*  
-- Ted Neward, 2004/6

*“O/RM is the Vietnam of Computer Science.”*  
-- Ted Neward, 2004/6

slippery slope

*“O/RM is the Vietnam of Computer Science.”*

-- Ted Neward, 2004/6

slippery slope

diminishing returns

# Impedance mismatch

# Impedance mismatch

- Relational = predicate logic, truth statements

# Impedance mismatch

- Relational = predicate logic, truth statements
- Object = identity, state, behavior, encapsulation

# Pick your poison

# Pick your poison

- Make the objects fit the database

# Pick your poison

- Make the objects fit the database
  - Active Record

# Pick your poison

- Make the objects fit the database
  - Active Record
- Make the database fit the objects

# Pick your poison

- Make the objects fit the database
  - Active Record
- Make the database fit the objects
  - OODBMS, Graph DBs

# Middle ground

# Middle ground

- Document databases (*identity and state*)

# Middle ground

- Document databases (*identity and state*)
  - No JOINs - denormalized / composed

# Middle ground

- Document databases (*identity and state*)
  - No JOINs - denormalized / composed
  - Loosely structured

# Middle ground

- Document databases (*identity and state*)
  - No JOINs - denormalized / composed
  - Loosely structured
  - Friendly transport formats, e.g. JSON

# What is Riak?

(the document database bits)

# What is Riak?

(the document database bits)

- Store your objects as JSON (or any format)

# What is Riak?

(the document database bits)

- Store your objects as JSON (or any format)
- Link between objects (like hypertext)

# What is Riak?

(the document database bits)

- Store your objects as JSON (or any format)
- Link between objects (like hypertext)
- No SQL - query with Javascript Map-Reduce

“DevOps” - agile infrastructure

“DevOps” - agile infrastructure  
programming your infrastructure

“DevOps” - agile infrastructure  
programming your infrastructure  
rapid “cloud” deployments

# Befriend your sysadmin

# Befriend your sysadmin

- Integrate with existing infrastructure

# Befriend your sysadmin

- Integrate with existing infrastructure
- Reduce the amount of hands-on work

# Befriend your sysadmin

- Integrate with existing infrastructure
- Reduce the amount of hands-on work
- Have a growth plan

# What is Riak?

(the ops-friendly bits)

# What is Riak?

(the ops-friendly bits)

- Web-shaped storage

# What is Riak?

(the ops-friendly bits)

- Web-shaped storage
  - Store data in its original format

# What is Riak?

(the ops-friendly bits)

- Web-shaped storage
  - Store data in its original format
  - It's just HTTP - same techniques apply

# What is Riak?

(the ops-friendly bits)

- Web-shaped storage
  - Store data in its original format
  - It's just HTTP - same techniques apply
    - load balancing, proxy caches, round-robin DNS

# What is Riak?

(the ops-friendly bits)

- Web-shaped storage
  - Store data in its original format
  - It's just HTTP - same techniques apply
    - load balancing, proxy caches, round-robin DNS
- No node is special - grow horizontally

# What is Riak?

(the ops-friendly bits)

- Web-shaped storage
  - Store data in its original format
  - It's just HTTP - same techniques apply
    - load balancing, proxy caches, round-robin DNS
- No node is special - grow horizontally
- Get some sleep, fix it in the morning

To review, Riak is...

# To review, Riak is...

- Magical Horizontally-Scaling Unicorns



# To review, Riak is...

- Magical Horizontally-Scaling Unicorns
- Application-Friendly Rainbows in the Cloud



# To review, Riak is...

- Magical Horizontally-Scaling Unicorns
- Application-Friendly Rainbows in the Cloud
- Data-Shredding MapReduce Ninjas



<http://downloads.basho.com/riak/>

hg clone <http://hg.basho.com/riak/>

gem install ripple

git clone git://github.com/seancribbs/ripple.git

# CRUD in Riak

# CRUD in Riak

- PUT /riak/bucket/key (C,U) (also POST)

# CRUD in Riak

- PUT /riak/bucket/key (C,U) (also POST)
- GET /riak/bucket/key (R)

# CRUD in Riak

- PUT /riak/bucket/key (C,U) (also POST)
- GET /riak/bucket/key (R)
- DELETE /riak/bucket/key (D)

# CRUD Demo

# Links

# Links

- One-way, qualified “pointers” to other objects

# Links

- One-way, qualified “pointers” to other objects
- “Link” header in HTTP

# Links

- One-way, qualified “pointers” to other objects
- “Link” header in HTTP
- </riak/bucket/key>; riaktag=”tag”

# Link-walking

# Link-walking

- GET /riak/bucket/key/[bucket,tag,keep]\*

# Link-walking

- GET /riak/bucket/key/[bucket,tag,keep]\*
  - bucket = scope to a bucket

# Link-walking

- GET /riak/bucket/key/[bucket,tag,keep]\*
  - bucket = scope to a bucket
  - tag = scope to a tag

# Link-walking

- GET /riak/bucket/key/[bucket,tag,keep]\*
  - bucket = scope to a bucket
  - tag = scope to a tag
  - keep = return results (last one always)

# Link-walking

- GET /riak/bucket/key/[bucket,tag,keep]\*
  - bucket = scope to a bucket
  - tag = scope to a tag
  - keep = return results (last one always)
  - “\_” means match anything

# Link-walking

- GET /riak/bucket/key/[bucket,tag,keep]\*
  - bucket = scope to a bucket
  - tag = scope to a tag
  - keep = return results (last one always)
  - “\_” means match anything
- Example:

# Link-walking

- GET /riak/bucket/key/[bucket,tag,keep]\*
  - bucket = scope to a bucket
  - tag = scope to a tag
  - keep = return results (last one always)
  - “\_” means match anything
- Example:
  - GET /riak/artists/TheBeatles/albums,\_,1/tracks,\_,1

# Links Demo

## Fault-tolerance



# Riak Map-Reduce

I believe I did, Bob.

# Riak Map-Reduce

# Riak Map-Reduce

- Based on Google's Map-Reduce idea

# Riak Map-Reduce

- Based on Google's Map-Reduce idea
- Some similarities to Hadoop and CouchDB

# Riak Map-Reduce

- Based on Google's Map-Reduce idea
- Some similarities to Hadoop and CouchDB
- High data-locality

# Riak Map-Reduce

- Based on Google's Map-Reduce idea
- Some similarities to Hadoop and CouchDB
- High data-locality
- Phases can be Javascript or Erlang

# Map-Reduce Terminology

# Map-Reduce Terminology

- **Job:** a query, composed of inputs and phases

# Map-Reduce Terminology

- **Job:** a query, composed of inputs and phases
- **Phase:** a single map or reduce function application

# Map-Reduce Terminology

- **Job:** a query, composed of inputs and phases
- **Phase:** a single map or reduce function application
- **Map:** phase applied to each item - filter, transform

# Map-Reduce Terminology

- **Job:** a query, composed of inputs and phases
- **Phase:** a single map or reduce function application
- **Map:** phase applied to each item - filter, transform
- **Reduce:** phase applied to the collection - collate, aggregate, compute

# Map Phases

# Map Phases

- Inputs are bucket-key pairs (with optional key-specific data)

# Map Phases

- Inputs are bucket-key pairs (with optional key-specific data)
- Must return a list

# Map Phases

- Inputs are bucket-key pairs (with optional key-specific data)
- Must return a list
- Parallel results are aggregated

# Map Built-ins

# Map Built-ins

- `Riak.mapValues`

# Map Built-ins

- `Riak.mapValues`
- `Riak.mapValuesJson`

# Reduce Phases

# Reduce Phases

- Performed on a single node

# Reduce Phases

- Performed on a single node
- Two processes per reduce phase increase parallelism

# Reduce Phases

- Performed on a single node
- Two processes per reduce phase increase parallelism
- Must return a list

# Reduce Built-ins

# Reduce Built-ins

- Riak. reduceMin

# Reduce Built-ins

- Riak.reduceMin
- Riak.reduceMax

# Reduce Built-ins

- Riak.reduceMin
- Riak.reduceMax
- Riak.reduceSort

# Build a M/R Job

# Build a M/R Job

- Specify inputs

# Build a M/R Job

- Specify inputs
  - bucket/key, bucket/key/key-specific-arg, bucket

# Build a M/R Job

- Specify inputs
  - bucket/key, bucket/key/key-specific-arg, bucket
- Specify phases

# Build a M/R Job

- Specify inputs
  - bucket/key, bucket/key/key-specific-arg, bucket
- Specify phases
  - Each can receive a static argument

# Build a M/R Job

- Specify inputs
  - bucket/key, bucket/key/key-specific-arg, bucket
- Specify phases
  - Each can receive a static argument
  - Each can return intermediate results

# Map-Reduce on HTTP

# Map-Reduce on HTTP

- Job is a JSON object

# Map-Reduce on HTTP

- Job is a JSON object
- POST /mapred

# A simple M/R job

```
{"inputs": "goog",
"query": [{"map": {
    "language": "javascript",
    "name": "Riak.mapValuesJson",
    "keep": true}}]}
```

# A simple M/R job

```
{"inputs": "goog",
"query": [{"map": {
    "language": "javascript",
    "name": "Riak.mapValuesJson",
    "keep": true}}]

Riak::MapReduce.new(c).add('goog').
  map('Riak.mapValuesJson', :keep => true).run
```

# Another M/R Job

```
{"inputs": "stocks",

"query": [{"map": {"language": "javascript",
            "name": "App.extractTickers",
            "arg": "GOOG",
            "keep": false},
          {"reduce": {"language": "javascript",
                     "name": "Riak.reduceMin",
                     "keep": true} }]}
```

# Another M/R Job

```
{"inputs": "stocks",

"query": [{"map": {"language": "javascript",
            "name": "App.extractTickers",
            "arg": "GOOG",
            "keep": false},
          {"reduce": {"language": "javascript",
                     "name": "Riak.reduceMin",
                     "keep": true}}]}
```

```
Riak::MapReduce.new(c).add('stocks').
  map('App.extractTickers', :arg =>"GOOG").
  reduce("Riak.reduceMin", :keep => true).run
```

# Map-Reduce Demo

# Ripple

# Ripple

- Document-style persistence (like MongoMapper)

# Ripple

- Document-style persistence (like MongoMapper)
- ActiveRecord niceties

# Ripple

- Document-style persistence (like MongoMapper)
- ActiveRecord niceties
- Callbacks

# Ripple

- Document-style persistence (like MongoMapper)
- ActiveRecord niceties
  - Callbacks
  - Validations

# Ripple

- Document-style persistence (like MongoMapper)
- ActiveRecord niceties
  - Callbacks
  - Validations
  - ActionPack compatibility

# Ripple TODO

# Ripple TODO

- Associations

# Ripple TODO

- Associations
- Indexes

# Ripple TODO

- **Associations**
- **Indexes**
- Use a property/method as the key

# Ripple TODO

- **Associations**
- **Indexes**
- Use a property/method as the key
- Dynamic finders

# Ripple TODO

- **Associations**
- **Indexes**
- Use a property/method as the key
- Dynamic finders
- Session/cache stores

# Ripple Demo

[sean@basho.com](mailto:sean@basho.com)

@seancribbs

“seancribbs” on Freenode IRC #riak

# Questions?

<http://wiki.basho.com>

[riak-users@lists.basho.com](mailto:riak-users@lists.basho.com)